

health-insurance-premium-analysis_v3

March 16, 2025

0.1 # Regression with an Insurance Dataset (Model Training)

Source: Kaggle Playground Prediction Competition

Course Title: DAMO-510-4: Winter 2025 Predictive Analytics

Professor Name: Professor Ali El-Sharif

Submission Date: March 9, 2025

Submitted By:

1. Denisse C. Cortes (NF1007936)
2. Dev D. Rabadia (NF1005560)
3. Miko L. Tan (NF1008647)
4. Rosario D. Torres (NF1001385)

The data model for this project is a regression-based model designed to predict insurance premium amounts. The model uses a variety of features, including numerical variables (such as age, annual income, and health score) and categorical variables (such as gender, marital status, and policy type) to estimate the target variable: the insurance premium amount. The dataset incorporates several features with skewed distributions and missing values, which will be handled through appropriate preprocessing techniques. Outliers in the dataset, such as those found in the “Previous Claims” feature, will also be addressed to improve model accuracy and robustness.

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import boxcox
from sklearn.feature_selection import RFE
# from sklearn.linear_model import LinearRegression
from statsmodels.stats.outliers_influence import variance_inflation_factor
from statsmodels.tools.tools import add_constant
import missingno as msno
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import LabelEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
```

```

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.tree import DecisionTreeRegressor
# from sklearn.ensemble import RandomForestRegressor
import xgboost as xgb
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import mean_squared_log_error
from sklearn.model_selection import ParameterGrid
from sklearn.linear_model import LinearRegression, Lasso, ElasticNet
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.svm import SVR
from xgboost import XGBRegressor

```

```

[2]: # Load the dataset
insurance_train_data = pd.read_csv('kaggle/input/playground-series-s4e12/train.
↳CSV')

```

0.2 Exploratory Data Analysis (EDA)

```

[3]: insurance_train_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1200000 entries, 0 to 1199999
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     1200000 non-null  int64
1   Age                                   1181295 non-null  float64
2   Gender                               1200000 non-null  object
3   Annual Income                         1155051 non-null  float64
4   Marital Status                        1181471 non-null  object
5   Number of Dependents                  1090328 non-null  float64
6   Education Level                       1200000 non-null  object
7   Occupation                            841925 non-null   object
8   Health Score                          1125924 non-null  float64
9   Location                              1200000 non-null  object
10  Policy Type                           1200000 non-null  object
11  Previous Claims                       835971 non-null   float64
12  Vehicle Age                           1199994 non-null  float64
13  Credit Score                          1062118 non-null  float64
14  Insurance Duration                    1199999 non-null  float64
15  Policy Start Date                     1200000 non-null  object
16  Customer Feedback                     1122176 non-null  object
17  Smoking Status                        1200000 non-null  object
18  Exercise Frequency                    1200000 non-null  object
19  Property Type                         1200000 non-null  object
20  Premium Amount                       1200000 non-null  float64

```

dtypes: float64(9), int64(1), object(11)
memory usage: 192.3+ MB

```
[4]: # Check for nulls
# insurance_train_data.isna().sum()

# Assuming 'insurance_train_data' is your dataset
missing_data = insurance_train_data.isnull().sum()

# Calculate percentage of missing data
missing_percentage = (missing_data / len(insurance_train_data)) * 100

# Create a dataframe to store the missing data information
missing_info = pd.DataFrame({
    'Missing Count': missing_data,
    'Missing Percentage': missing_percentage,
})

# Define categories for missing percentage
def categorize_missing_data(percentage):
    if percentage <= 5:
        return 'Small (1-5%)'
    elif 5 < percentage <= 20:
        return 'Moderate (5-20%)'
    elif 20 < percentage <= 40:
        return 'High (20-40%)'
    else:
        return 'Very High (40%+)

# Apply the categorization function
missing_info['Classification'] = missing_info['Missing Percentage'].
    ↪apply(categorize_missing_data)

# Sort by missing percentage in descending order for better visibility
missing_info = missing_info.sort_values(by='Missing Percentage',
    ↪ascending=False)

# Display the result
print(missing_info)
```

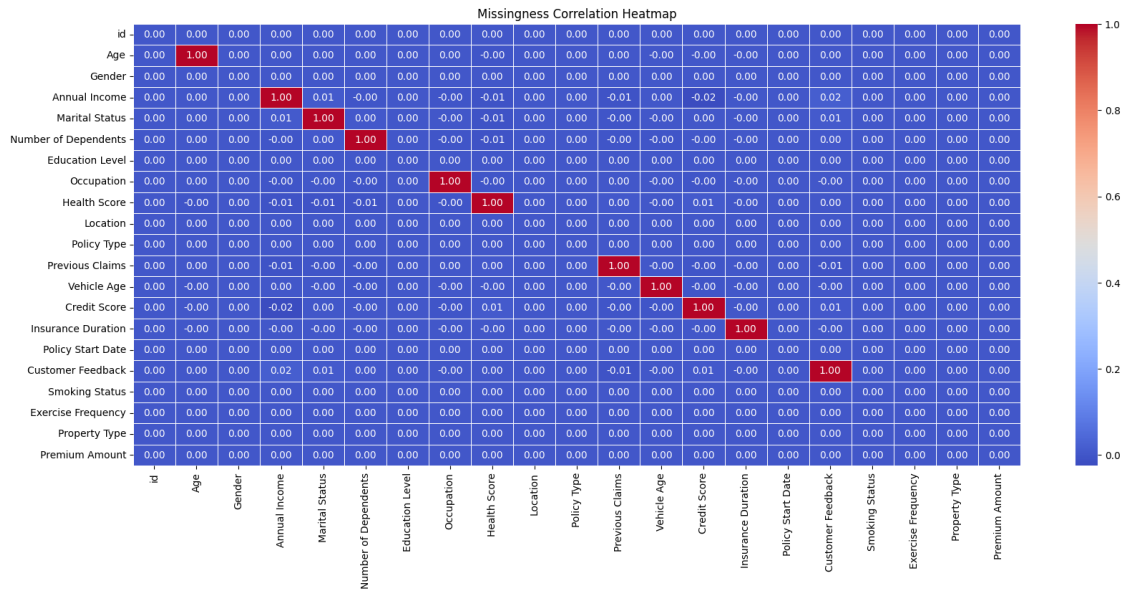
	Missing Count	Missing Percentage	Classification
Previous Claims	364029	30.335750	High (20-40%)
Occupation	358075	29.839583	High (20-40%)
Credit Score	137882	11.490167	Moderate (5-20%)
Number of Dependents	109672	9.139333	Moderate (5-20%)
Customer Feedback	77824	6.485333	Moderate (5-20%)
Health Score	74076	6.173000	Moderate (5-20%)
Annual Income	44949	3.745750	Small (1-5%)

Age	18705	1.558750	Small (1-5%)
Marital Status	18529	1.544083	Small (1-5%)
Vehicle Age	6	0.000500	Small (1-5%)
Insurance Duration	1	0.000083	Small (1-5%)
Gender	0	0.000000	Small (1-5%)
id	0	0.000000	Small (1-5%)
Location	0	0.000000	Small (1-5%)
Policy Type	0	0.000000	Small (1-5%)
Education Level	0	0.000000	Small (1-5%)
Policy Start Date	0	0.000000	Small (1-5%)
Smoking Status	0	0.000000	Small (1-5%)
Exercise Frequency	0	0.000000	Small (1-5%)
Property Type	0	0.000000	Small (1-5%)
Premium Amount	0	0.000000	Small (1-5%)

```
[5]: # Missingness Analysis (v2)
# Identify which variables tend to have missing values at the same time
# Create a DataFrame with missingness indicator (1 for missing, 0 for not
↳missing)
missing_data = insurance_train_data.isnull().astype(int)

# Calculate the correlation matrix of the missingness
corr_matrix = missing_data.corr()
# Fill NaN values before plotting
corr_matrix = corr_matrix.fillna(0)
plt.figure(figsize=(20, 8))

# Plot the correlation matrix as a heatmap
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", linewidths=0.5, fmt='.2f')
plt.title("Missingness Correlation Heatmap")
plt.show()
```



```
[6]: # Check for duplicate rows
insurance_train_data.duplicated().sum()
```

```
[6]: np.int64(0)
```

```
[7]: insurance_train_data.nunique()
```

```
[7]: id          1200000
Age             47
Gender           2
Annual Income   88593
Marital Status   3
Number of Dependents  5
Education Level  4
Occupation       3
Health Score     532657
Location         3
Policy Type      3
Previous Claims  10
Vehicle Age     20
Credit Score     550
Insurance Duration  9
Policy Start Date 167381
Customer Feedback  3
Smoking Status   2
Exercise Frequency  4
Property Type     3
Premium Amount   4794
```

dtype: int64

```
[8]: # Describe the dataset
# insurance_train_data.describe()

# Apply styling to the describe output
styled_description = insurance_train_data.describe(include='all').style \
    .format({
        'Age': '{:.0f}',
        'Annual Income': '{:.2f}',
        'Number of Dependents': '{:.0f}',
        'Health Score': '{:.0f}',
        'Previous Claims': '{:.2f}',
        'Vehicle Age': '{:.0f}',
        'Credit Score': '{:.0f}',
        'Insurance Duration': '{:.2f}',
        'Premium Amount': '{:.2f}',
    }) \
    .map(lambda x: 'color: green' if isinstance(x, (int, float)) and x < 0 else
    ↪ '', subset=['Premium Amount'])

# Display the styled description
styled_description
```

[8]: <pandas.io.formats.style.Styler at 0x262000cf550>

```
[9]: # Check sample values
insurance_train_data.head()
```

```
[9]:
```

	id	Age	Gender	Annual Income	Marital Status	Number of Dependents	\
0	0	19.0	Female	10049.0	Married	1.0	
1	1	39.0	Female	31678.0	Divorced	3.0	
2	2	23.0	Male	25602.0	Divorced	3.0	
3	3	21.0	Male	141855.0	Married	2.0	
4	4	21.0	Male	39651.0	Single	1.0	

	Education Level	Occupation	Health Score	Location	...	Previous Claims	\
0	Bachelor's	Self-Employed	22.598761	Urban	...	2.0	
1	Master's	NaN	15.569731	Rural	...	1.0	
2	High School	Self-Employed	47.177549	Suburban	...	1.0	
3	Bachelor's	NaN	10.938144	Rural	...	1.0	
4	Bachelor's	Self-Employed	20.376094	Rural	...	0.0	

	Vehicle Age	Credit Score	Insurance Duration	Policy Start Date	\
0	17.0	372.0	5.0	2023-12-23 15:21:39.134960	
1	12.0	694.0	2.0	2023-06-12 15:21:39.111551	
2	14.0	NaN	3.0	2023-09-30 15:21:39.221386	

3	0.0	367.0	1.0	2024-06-12 15:21:39.226954
4	8.0	598.0	4.0	2021-12-01 15:21:39.252145

	Customer Feedback	Smoking Status	Exercise Frequency	Property Type \
0	Poor	No	Weekly	House
1	Average	Yes	Monthly	House
2	Good	Yes	Weekly	House
3	Poor	Yes	Daily	Apartment
4	Poor	Yes	Weekly	House

	Premium Amount
0	2869.0
1	1483.0
2	567.0
3	765.0
4	2022.0

[5 rows x 21 columns]

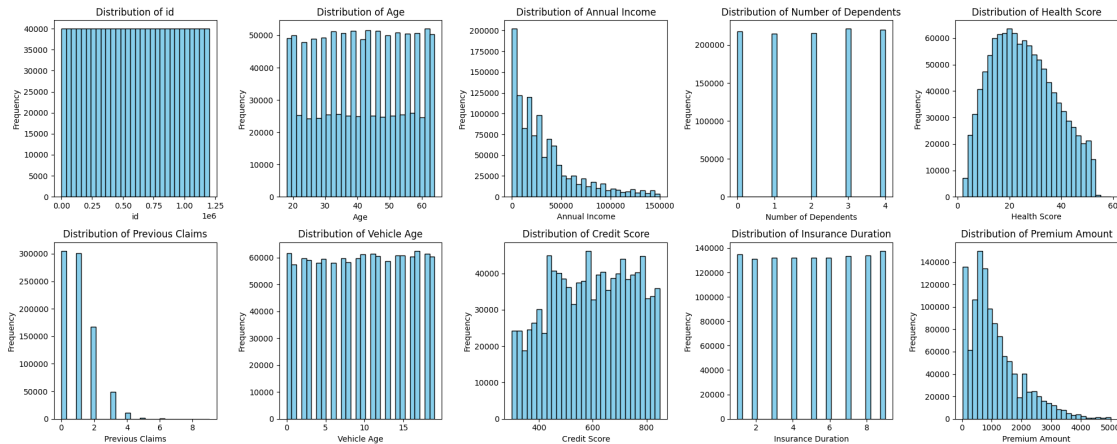
```
[10]: # Check distribution of numerical features
# Get numerical columns
num_columns = insurance_train_data.select_dtypes(include=['number']).columns

# Define grid size (5 columns)
num_cols = 5
num_rows = (len(num_columns) // num_cols) + (len(num_columns) % num_cols > 0)

# Create a figure with a grid of subplots
plt.figure(figsize=(20, 4 * num_rows)) # Adjust height based on the number of
    ↪ rows

# Loop through numerical columns and plot histograms
for i, col in enumerate(num_columns, 1):
    plt.subplot(num_rows, num_cols, i) # Position plot in a 4x5 grid
    plt.hist(insurance_train_data[col], bins=30, color='skyblue',
    ↪ edgecolor='black')
    plt.title(f"Distribution of {col}")
    plt.xlabel(col)
    plt.ylabel("Frequency")

# Adjust layout to prevent overlap
plt.tight_layout()
plt.show()
```



```
[11]: # Skewness Computation with Interpretation
# Select only numerical columns
numerical_columns = insurance_train_data.select_dtypes(include=['number'])

# Compute skewness
skew_values = numerical_columns.skew().sort_values(ascending=False)

# Create a DataFrame to store skewness values
skew_df = pd.DataFrame({'Skewness Value': skew_values})

# Apply interpretation directly
skew_df['Interpretation'] = ""

skew_df.loc[skew_df['Skewness Value'] == 0, 'Interpretation'] = "Norm Dist ( = 0 )"
skew_df.loc[(skew_df['Skewness Value'] > -0.5) & (skew_df['Skewness Value'] < 0.5), 'Interpretation'] = "Min/No Skew ( = -0.5 to 0.5 )"
skew_df.loc[skew_df['Skewness Value'] > 1, 'Interpretation'] = "Extreme Right-Skewed ( > 1 )"
skew_df.loc[skew_df['Skewness Value'] < -1, 'Interpretation'] = "Extreme Left-Skewed ( < -1 )"
skew_df.loc[(skew_df['Skewness Value'] >= 0.5) & (skew_df['Skewness Value'] <= 1), 'Interpretation'] = "Right-Skewed ( > 0.5 )"
skew_df.loc[(skew_df['Skewness Value'] <= -0.5) & (skew_df['Skewness Value'] >= -1), 'Interpretation'] = "Left-Skewed ( < -0.5 )"

# Display the skewness table
print(skew_df)
```

	Skewness Value	Interpretation
Annual Income	1.470357e+00	Extreme Right-Skewed (> 1)
Premium Amount	1.240915e+00	Extreme Right-Skewed (> 1)

Previous Claims	9.053210e-01	Right-Skewed (> 0.5)
Health Score	2.821873e-01	Min/No Skew (= -0.5 to 0.5)
id	3.836279e-16	Min/No Skew (= -0.5 to 0.5)
Insurance Duration	-8.793302e-03	Min/No Skew (= -0.5 to 0.5)
Age	-1.253192e-02	Min/No Skew (= -0.5 to 0.5)
Number of Dependents	-1.325461e-02	Min/No Skew (= -0.5 to 0.5)
Vehicle Age	-2.040888e-02	Min/No Skew (= -0.5 to 0.5)
Credit Score	-1.135726e-01	Min/No Skew (= -0.5 to 0.5)

```
[12]: # Boxplots for outlier detection

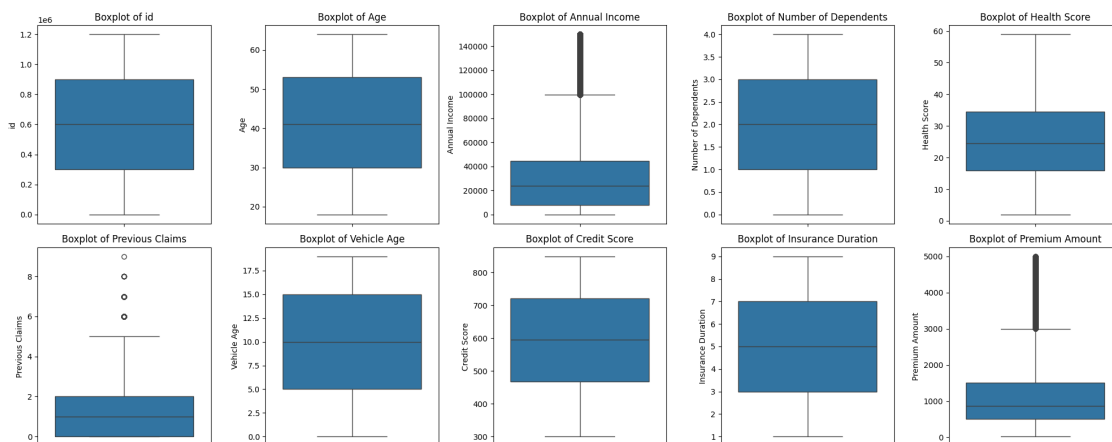
# Get numerical columns
num_columns = insurance_train_data.select_dtypes(include=['number']).columns

# Define grid size (5 columns)
num_cols = 5
num_rows = (len(num_columns) // num_cols) + (len(num_columns) % num_cols > 0)

# Create a figure with a grid of subplots
plt.figure(figsize=(20, 4 * num_rows)) # Adjust height based on the number of
↪rows

# Loop through numerical columns and create boxplots
for i, col in enumerate(num_columns, 1):
    plt.subplot(num_rows, num_cols, i) # Position plot in a 4x5 grid
    sns.boxplot(y=insurance_train_data[col])
    plt.title(f"Boxplot of {col}")

# Adjust layout to prevent overlap
plt.tight_layout()
plt.show()
```



```
[13]: # Print unique values for categorical columns
for col in insurance_train_data.select_dtypes(include=["object"]).columns:
    unique_vals = insurance_train_data[col].dropna().unique()
    print(f"Column: {col}")
    print("Unique Values:", unique_vals)
    print("Number of Unique Values:", len(unique_vals))
    print("-" * 50)
```

Column: Gender

Unique Values: ['Female' 'Male']

Number of Unique Values: 2

Column: Marital Status

Unique Values: ['Married' 'Divorced' 'Single']

Number of Unique Values: 3

Column: Education Level

Unique Values: ["Bachelor's" "Master's" 'High School' 'PhD']

Number of Unique Values: 4

Column: Occupation

Unique Values: ['Self-Employed' 'Employed' 'Unemployed']

Number of Unique Values: 3

Column: Location

Unique Values: ['Urban' 'Rural' 'Suburban']

Number of Unique Values: 3

Column: Policy Type

Unique Values: ['Premium' 'Comprehensive' 'Basic']

Number of Unique Values: 3

Column: Policy Start Date

Unique Values: ['2023-12-23 15:21:39.134960' '2023-06-12 15:21:39.111551'

'2023-09-30 15:21:39.221386' ... '2021-04-28 15:21:39.129190'

'2019-11-14 15:21:39.201446' '2020-10-19 15:21:39.118178']

Number of Unique Values: 167381

Column: Customer Feedback

Unique Values: ['Poor' 'Average' 'Good']

Number of Unique Values: 3

Column: Smoking Status

Unique Values: ['No' 'Yes']

Number of Unique Values: 2

Column: Exercise Frequency

Unique Values: ['Weekly' 'Monthly' 'Daily' 'Rarely']

Number of Unique Values: 4

Column: Property Type

Unique Values: ['House' 'Apartment' 'Condo']

Number of Unique Values: 3

```
[14]: # Get categorical columns (v2)
cat_columns = insurance_train_data.select_dtypes(include=["object"]).columns

# Exclude 'Policy Start Date' from the list of categorical columns
cat_columns = [col for col in cat_columns if col != 'Policy Start Date']

# Define the grid size (5 columns)
num_cols = 5
num_rows = (len(cat_columns) // num_cols) + (len(cat_columns) % num_cols > 0)

# Create a figure with a grid of subplots
plt.figure(figsize=(20, 4 * num_rows)) # Adjust height based on the number of
    ↪ rows

# Loop through categorical columns and plot their distributions
for i, col in enumerate(cat_columns, 1):
    plt.subplot(num_rows, num_cols, i) # Position plot in 4x5 grid

    ax = sns.countplot(x=insurance_train_data[col],
    ↪ order=insurance_train_data[col].value_counts().index,
    hue=insurance_train_data[col], palette="viridis") #
    ↪ , legend=False

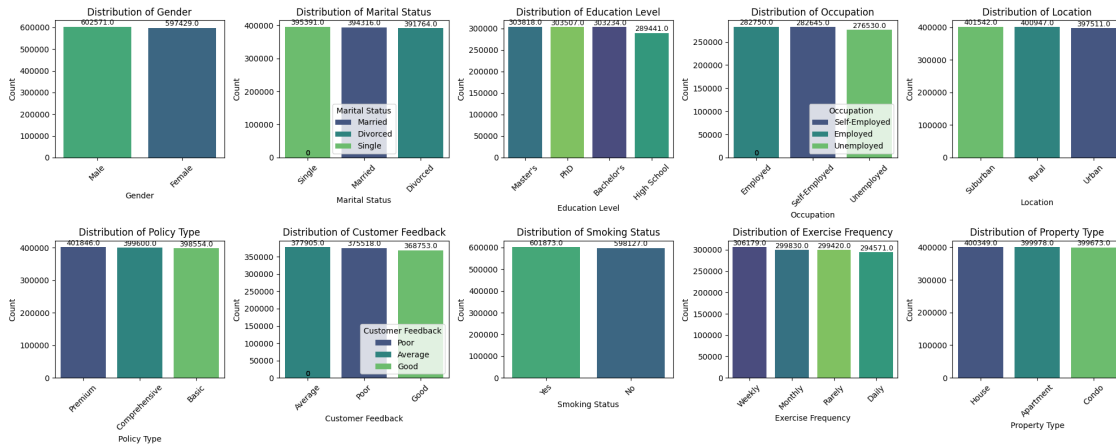
    # Add data labels to the bars
    for p in ax.patches:
        ax.annotate(f'{p.get_height()}',
                    (p.get_x() + p.get_width() / 2., p.get_height()),
                    ha='center', va='center',
                    fontsize=9, color='black',
                    xytext=(0, 5), textcoords='offset points')

    # Rotate x-ticks for better readability
    plt.xticks(rotation=45)

    # Title and labels
    plt.title(f"Distribution of {col}")
    plt.xlabel(col)
    plt.ylabel("Count")

# Adjust layout to prevent overlap
```

```
plt.tight_layout()
plt.show()
```



```
[15]: # Display category counts instead of plotting
for col in cat_columns:
    print(f"\n{col} Distribution:\n", insurance_train_data[col].value_counts())
```

Gender Distribution:

Gender

Male 602571

Female 597429

Name: count, dtype: int64

Marital Status Distribution:

Marital Status

Single 395391

Married 394316

Divorced 391764

Name: count, dtype: int64

Education Level Distribution:

Education Level

Master's 303818

PhD 303507

Bachelor's 303234

High School 289441

Name: count, dtype: int64

Occupation Distribution:

Occupation

Employed 282750

Self-Employed 282645
Unemployed 276530
Name: count, dtype: int64

Location Distribution:

Location
Suburban 401542
Rural 400947
Urban 397511
Name: count, dtype: int64

Policy Type Distribution:

Policy Type
Premium 401846
Comprehensive 399600
Basic 398554
Name: count, dtype: int64

Customer Feedback Distribution:

Customer Feedback
Average 377905
Poor 375518
Good 368753
Name: count, dtype: int64

Smoking Status Distribution:

Smoking Status
Yes 601873
No 598127
Name: count, dtype: int64

Exercise Frequency Distribution:

Exercise Frequency
Weekly 306179
Monthly 299830
Rarely 299420
Daily 294571
Name: count, dtype: int64


Property Type Distribution:

Property Type
House 400349
Apartment 399978
Condo 399673
Name: count, dtype: int64

```
[16]: # Select only numerical columns
num_features = insurance_train_data.select_dtypes(include=['number'])

# Compute correlation matrix
corr_matrix = num_features.corr()

# Plot heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", fmt=".2f", linewidths=0.5)
plt.title("Correlation Heatmap (Numerical Features Only)")
plt.show()
```




```
[17]: # v2
# Get correlation with target variable "Premium Amount"
corr_with_target = corr_matrix["Premium Amount"].drop("Premium Amount")

# Visualizing Correlation with Target Variable
plt.figure(figsize=(6, 8))
sns.barplot(
```

```

y=corr_with_target.sort_values(ascending=False).index,
x=corr_with_target.sort_values(ascending=False).values,
hue=corr_with_target.sort_values(ascending=False).index, # Assign y to hue
palette="coolwarm"
# legend=False # Hide legend since hue is only for coloring
)
plt.title("Correlation with Premium Amount (Numerical Features Only)")
plt.xlabel("Correlation Score")
plt.ylabel("Features")
plt.show()

```



0.3 Data Preprocessing

Handling Missing Values

```
[18]: # 1. Id - Not needed for analysis
# if "id" in insurance_train_data.columns:
#     insurance_train_data.drop(columns=["id"], inplace=True)
#     print("Column 'id' dropped successfully.")
# else:
#     print("Column 'id' does not exist in the dataset.")
# DO NOT DROP, needed for Submission File!

[19]: # 2. Age - Okay to drop Age if Small percentage of missing values, otherwise
      ↪impute
if missing_info.loc["Age"]["Missing Percentage"] < 5:
    insurance_train_data.dropna(subset=["Age"], inplace=True)
else:
    insurance_train_data["Age"].fillna(insurance_train_data["Age"].median(),
    ↪inplace=True)

[20]: # 3. Gender - Tag missing values as Unknown
insurance_train_data.loc[:, "Gender"] = insurance_train_data["Gender"].
    ↪fillna("Unknown")

[21]: # 4. Annual Income - Okay to drop Age if Small percentage of missing values,
      ↪otherwise impute
if missing_info.loc["Annual Income"]["Missing Percentage"] < 5:
    insurance_train_data.dropna(subset=["Annual Income"], inplace=True)
else:
    insurance_train_data["Annual Income"].fillna(insurance_train_data["Annual
    ↪Income"].median(), inplace=True)

[22]: # 5. Marital Status - Tag missing values as Single
# insurance_train_data["Marital Status"].replace({"Unknown": "Single"},
    ↪inplace=True)
insurance_train_data.loc[:, "Marital Status"] = insurance_train_data["Marital
    ↪Status"].fillna("Single")

[23]: # 6. Number of Dependents - Default missing values to 0 dependents
insurance_train_data.loc[:, "Number of Dependents"] =
    ↪insurance_train_data["Number of Dependents"].fillna(0)

[24]: # 7. Education Level - No missing values now, but in the future, default it to
      ↪Unknown
insurance_train_data.loc[:, "Education Level"] =
    ↪insurance_train_data["Education Level"].fillna("Unknown")
```



```

[25]: # 8. Occupation - Default to Unemployed
insurance_train_data.loc[:, "Occupation"] = insurance_train_data["Occupation"].
↳fillna("Unemployed")

[26]: # 9. Health Score - Default to 0
insurance_train_data.loc[:, "Health Score"] = insurance_train_data["Health_
↳Score"].fillna(0)

[27]: # 10. Location - No missing values now, but in the future, default it to Unknown
insurance_train_data.loc[:, "Location"] = insurance_train_data["Location"].
↳fillna("Unknown")

[28]: # 11. Policy Type - No missing values now, but in the future, default it to
↳Basic
insurance_train_data.loc[:, "Policy Type"] = insurance_train_data["Policy_
↳Type"].fillna("Basic")

[29]: # 12. Previous Claims - Default to 0, no claim
insurance_train_data.loc[:, "Previous Claims"] = insurance_train_data["Previous_
↳Claims"].fillna(0)

[30]: # 13. Vehicle Age - Default to 0, no vehicle declared
insurance_train_data.loc[:, "Vehicle Age"] = insurance_train_data["Vehicle_
↳Age"].fillna(0)

[31]: # 14. Credit Score - Default to 0, possibly newcomers
insurance_train_data.loc[:, "Credit Score"] = insurance_train_data["Credit_
↳Score"].fillna(0)

[32]: # 15. Insurance Duration - Default to 0, new policy holders
insurance_train_data.loc[:, "Insurance Duration"] =
↳insurance_train_data["Insurance Duration"].fillna(0)

[33]: # 16. Policy Start Date - No missing values now, but dropped if not provided
insurance_train_data.dropna(subset=["Policy Start Date"], inplace=True)

[34]: # 17. Customer Feedback - Default to Not Provided
insurance_train_data.loc[:, "Customer Feedback"] =
↳insurance_train_data["Customer Feedback"].fillna("Not Provided")

[35]: # 18. Smoking Status - No missing values now, but in the future, default to No
insurance_train_data.loc[:, "Smoking Status"] = insurance_train_data["Smoking_
↳Status"].fillna("No")

[36]: # 19. Exercise Frequency - No missing values now, but in the future, default to
↳Not Provided

```

```
insurance_train_data.loc[:, "Exercise Frequency"] =  
    ↪insurance_train_data["Exercise Frequency"].fillna("Not Provided")
```

```
[37]: # 20. Property Type - No missing values now, but in the future, default to Not  
    ↪Provided  
insurance_train_data.loc[:, "Property Type"] = insurance_train_data["Property_  
    ↪Type"].fillna("Not Provided")
```

```
[38]: # 21. Premium Amount - No missing values now, but in the future, drop if not  
    ↪provided  
insurance_train_data.dropna(subset=["Premium Amount"], inplace=True)
```

Feature Engineering & Transformation

```
[39]: # 1. Id - No further action
```

```
[40]: # 2. Age - Binning  
# Define the bins for age groups relevant for health insurance coverage  
bins = [0, 5, 18, 25, 35, 45, 55, 65, float('inf')] # Age bin edges  
labels = ["High Risk (0-5) (Infants)", "Moderate Risk (6-18) (Children &  
    ↪Adolescents)", "Low Risk (19-25) (Young Adults)",  
          "Moderate-Low Risk (26-35) (Early Adulthood)", "Moderate Risk (36-45)  
    ↪(Middle Adulthood)",  
          "High Risk (46-55) (Mature Adults)", "Very High Risk (56-65)  
    ↪(Pre-Retirement)", "Very High Risk (65+) (Seniors)"]  
  
# Create a new column with binned age groups  
# right=True includes the rightmost edge  
insurance_train_data["Age_Bin"] = pd.cut(insurance_train_data["Age"],  
    ↪bins=bins, labels=labels, right=True)  
# Check the result  
print(insurance_train_data['Age_Bin'].value_counts())  
  
# Define the mapping of age bin labels to numeric values  
age_bin_mapping = {  
    "High Risk (0-5) (Infants)": 1,  
    "Moderate Risk (6-18) (Children & Adolescents)": 2,  
    "Low Risk (19-25) (Young Adults)": 3,  
    "Moderate-Low Risk (26-35) (Early Adulthood)": 4,  
    "Moderate Risk (36-45) (Middle Adulthood)": 5,  
    "High Risk (46-55) (Mature Adults)": 6,  
    "Very High Risk (56-65) (Pre-Retirement)": 7,  
    "Very High Risk (65+) (Seniors)": 8  
}  
# Apply the mapping to the Age Bin column to convert it to numeric values  
insurance_train_data["Age_Bin Numeric"] = insurance_train_data["Age_Bin"].  
    ↪map(age_bin_mapping)
```

```
# Convert the binned column to integer codes
insurance_train_data["Age_Bin Numeric"] = insurance_train_data["Age_Bin_
↳Numeric"].cat.codes
# Check the result
print(insurance_train_data[['Age', 'Age_Bin', 'Age_Bin Numeric']].head())
```

```
Age_Bin
High Risk (46-55) (Mature Adults)          243558
Moderate Risk (36-45) (Middle Adulthood)   242815
Moderate-Low Risk (26-35) (Early Adulthood) 241035
Very High Risk (56-65) (Pre-Retirement)   220555
Low Risk (19-25) (Young Adults)            165605
Moderate Risk (6-18) (Children & Adolescents) 23518
High Risk (0-5) (Infants)                  0
Very High Risk (65+) (Seniors)             0
Name: count, dtype: int64
```

	Age	Age_Bin	Age_Bin Numeric
0	19.0	Low Risk (19-25) (Young Adults)	2
1	39.0	Moderate Risk (36-45) (Middle Adulthood)	4
2	23.0	Low Risk (19-25) (Young Adults)	2
3	21.0	Low Risk (19-25) (Young Adults)	2
4	21.0	Low Risk (19-25) (Young Adults)	2

```
[41]: # 3. Gender - One-Hot Encoding
# Check if the "Gender" column still exists in the DataFrame
if "Gender" in insurance_train_data.columns:
    # Perform one-hot encoding for "Gender" if it hasn't been done already
    insurance_train_data = pd.get_dummies(insurance_train_data,
↳columns=["Gender"], drop_first=False)
    print("Gender column has been one-hot encoded.")
else:
    print("Gender column has already been one-hot encoded.")

# Convert the boolean True/False values to 1/0
if 'Gender_Male' in insurance_train_data.columns:
    insurance_train_data['Gender_Male'] = insurance_train_data['Gender_Male'].
↳astype(int)
if 'Gender_Female' in insurance_train_data.columns:
    insurance_train_data['Gender_Female'] =
↳insurance_train_data['Gender_Female'].astype(int)
if 'Gender_Unknown' in insurance_train_data.columns:
    insurance_train_data['Gender_Unknown'] =
↳insurance_train_data['Gender_Unknown'].astype(int)
```

Gender column has been one-hot encoded.

```
[42]: # 4. Annual Income - Binning
# No need to perform Box-Cox Transformation as Binning was applied
# Define the income bins and labels
income_bins = [0, 25000, 50000, 100000, 150000, float('inf')] # Define bin
↳ edges
income_labels = ["Low Income (0-25k)", "Lower-Middle Income (25k-50k)",
                 "Middle Income (50k-100k)", "Upper-Middle Income (100k-150k)",
                 "High Income (150k+)"]

# Create a new column for binned income groups
# right=True includes the rightmost edge
insurance_train_data["Income_Bin"] = pd.cut(insurance_train_data["Annual_Income"], bins=income_bins, labels=income_labels, right=False)
# Check the result
print(insurance_train_data['Income_Bin'].value_counts())

# Define the mapping of income bin labels to numeric values
income_bin_mapping = {
    "Low Income (0-25k)": 1,
    "Lower-Middle Income (25k-50k)": 2,
    "Middle Income (50k-100k)": 3,
    "Upper-Middle Income (100k-150k)": 4,
    "High Income (150k+)": 5
}
# Apply the mapping to the Income Bin column to convert it to numeric values
insurance_train_data["Income_Bin Numeric"] = insurance_train_data["Income_Bin"].map(income_bin_mapping)
# Convert the binned column to integer codes
insurance_train_data["Income_Bin Numeric"] = insurance_train_data["Income_Bin Numeric"].cat.codes
# Check the result
print(insurance_train_data[['Annual Income', 'Income_Bin', 'Income_Bin Numeric']].head())
```

```
Income_Bin
Low Income (0-25k)          591399
Lower-Middle Income (25k-50k) 309977
Middle Income (50k-100k)    170188
Upper-Middle Income (100k-150k) 65522
High Income (150k+)         0
Name: count, dtype: int64
```

	Annual Income	Income_Bin	Income_Bin Numeric
0	10049.0	Low Income (0-25k)	0
1	31678.0	Lower-Middle Income (25k-50k)	1
2	25602.0	Lower-Middle Income (25k-50k)	1
3	141855.0	Upper-Middle Income (100k-150k)	3
4	39651.0	Lower-Middle Income (25k-50k)	1

```
[43]: # 5. Marital Status - One-Hot Encoding
# Check if the "Marital Status" column still exists in the DataFrame
if "Marital Status" in insurance_train_data.columns:
    # Perform one-hot encoding for "Marital Status" if it hasn't been done
    ↪already
    insurance_train_data = pd.get_dummies(insurance_train_data,
    ↪columns=["Marital Status"], drop_first=False)
    print("Marital Status column has been one-hot encoded.")
else:
    print("Marital Status column has already been one-hot encoded.")
```

Marital Status column has been one-hot encoded.

```
[44]: # 6. Number of Dependents - No action required.
```

```
[45]: # 7. Education Level - Label Encoding because of sequence
# Initialize the LabelEncoder
label_encoder = LabelEncoder()

# Define the desired order for the 'Education Level' categories
desired_order = ["Unknown", "High School", "Bachelor's", "Master's", "PhD"]

# Ensure the encoder fits based on the custom order
label_encoder.classes_ = np.array(desired_order)

# Apply the label encoder
insurance_train_data['Education Level_Encoded'] = label_encoder.
    ↪transform(insurance_train_data['Education Level'])
# Check the result
print(insurance_train_data['Education Level_Encoded'].value_counts())
```

```
Education Level_Encoded
3    287723
4    287653
2    287337
1    274373
Name: count, dtype: int64
```

```
[46]: # 8. Occupation - One-Hot Encoding
# Check if the "Occupation" column still exists in the DataFrame
if "Occupation" in insurance_train_data.columns:
    # Merge Employed and Self-Employed into 'Employed/Self-Employed', and keep
    ↪'Unemployed' as it is
    insurance_train_data['Occupation'] = insurance_train_data['Occupation'].
    ↪replace({
        'Employed': 'Employed/Self-Employed',
        'Self-Employed': 'Employed/Self-Employed'
    })
```

```

    # One-hot encode the 'Occupation' column
    insurance_train_data = pd.get_dummies(insurance_train_data,
    ↪columns=['Occupation'])
    print("Occupation column has been merged and one-hot encoded.")
else:
    print("Occupation column has already been merged and one-hot encoded.")

# Convert the boolean True/False values to 1/0
if 'Occupation_Employed/Self-Employed' in insurance_train_data.columns:
    insurance_train_data['Occupation_Employed/Self-Employed'] =
    ↪insurance_train_data['Occupation_Employed/Self-Employed'].astype(int)
if 'Occupation_Unemployed' in insurance_train_data.columns:
    insurance_train_data['Occupation_Unemployed'] =
    ↪insurance_train_data['Occupation_Unemployed'].astype(int)

```

Occupation column has been merged and one-hot encoded.

[47]: # 9. Health Score - No action required.

```

[48]: # 10. Location - One-Hot Encoding
# Check if the "Location" column still exists in the DataFrame
if "Location" in insurance_train_data.columns:
    # Perform one-hot encoding for "Location" if it hasn't been done already
    insurance_train_data = pd.get_dummies(insurance_train_data,
    ↪columns=["Location"], drop_first=False)
    print("Location column has been one-hot encoded.")
else:
    print("Location column has already been one-hot encoded.")

# Convert the boolean True/False values to 1/0
if 'Location_Suburban' in insurance_train_data.columns:
    insurance_train_data['Location_Suburban'] =
    ↪insurance_train_data['Location_Suburban'].astype(int)
if 'Location_Rural' in insurance_train_data.columns:
    insurance_train_data['Location_Rural'] =
    ↪insurance_train_data['Location_Rural'].astype(int)
if 'Location_Urban' in insurance_train_data.columns:
    insurance_train_data['Location_Urban'] =
    ↪insurance_train_data['Location_Urban'].astype(int)
if 'Location_Unknown' in insurance_train_data.columns:
    insurance_train_data['Location_Unknown'] =
    ↪insurance_train_data['Location_Unknown'].astype(int)

```

Location column has been one-hot encoded.

```

[49]: # 11. Policy Type - One-Hot Encoding
# Check if the "Policy Type" column still exists in the DataFrame
if "Policy Type" in insurance_train_data.columns:

```

```

    # Perform one-hot encoding for "Policy Type" if it hasn't been done already
    insurance_train_data = pd.get_dummies(insurance_train_data,
columns=["Policy Type"], drop_first=False)
    print("Policy Type column has been one-hot encoded.")
else:
    print("Policy Type column has already been one-hot encoded.")

# Convert the boolean True/False values to 1/0
if 'Policy Type_Basic' in insurance_train_data.columns:
    insurance_train_data['Policy Type_Basic'] = insurance_train_data['Policy_
Type_Basic'].astype(int)
if 'Policy Type_Comprehensive' in insurance_train_data.columns:
    insurance_train_data['Policy Type_Comprehensive'] =
insurance_train_data['Policy Type_Comprehensive'].astype(int)
if 'Policy Type_Premium' in insurance_train_data.columns:
    insurance_train_data['Policy Type_Premium'] = insurance_train_data['Policy_
Type_Premium'].astype(int)

```

Policy Type column has been one-hot encoded.

```

[50]: # 12. Previous Claims - Apply Box-Cox Transformation
# Select skewed numerical columns
skewed_features = ['Previous Claims']

# Ensure all values are positive (Box-Cox requires positive numbers)
insurance_train_data[skewed_features] = insurance_train_data[skewed_features].
    apply(lambda x: x + 1 if (x <= 0).any() else x)

# Apply Box-Cox transformation
for col in skewed_features:
    insurance_train_data[col], _ = boxcox(insurance_train_data[col])
# Check the result
print(insurance_train_data['Previous Claims'].value_counts())

```

```

Previous Claims
0.000000    635423
0.468220    285580
0.605892    158099
0.669637     45764
0.705798      9967
0.728860     1904
0.744740       285
0.756282        55
0.765017         8
0.771838         1
Name: count, dtype: int64

```

```

[51]: # 13. Vehicle Age - No action required.

```

```

[52]: # 14. Credit Score - No action required.

[53]: # 15. Insurance Duration - No action required.

[54]: # 16. Policy Start Date - Extract Temporal Patters.
# Convert the 'Policy Start Date' to datetime format
insurance_train_data['Policy Start Date'] = pd.
    ↪to_datetime(insurance_train_data['Policy Start Date'], errors='coerce')

# Extract year from the date
insurance_train_data['Policy Start Year'] = insurance_train_data['Policy Start_
    ↪Date'].dt.year

# Extract the month from the date
insurance_train_data['Policy Start Month'] = insurance_train_data['Policy Start_
    ↪Date'].dt.month

# Extract the difference in years from the current date
insurance_train_data['Years Since Start'] = (pd.to_datetime('today') -
    ↪insurance_train_data['Policy Start Date']).dt.days / 365

# Drop the original 'Policy Start Date' if no longer needed
insurance_train_data.drop(columns=["Policy Start Date"], inplace=True)

[55]: # 17. Customer Feedback - One-Hot Encoding
# Check if the "Customer Feedback" column still exists in the DataFrame
if "Customer Feedback" in insurance_train_data.columns:
    # Perform one-hot encoding for "Customer Feedback" if it hasn't been done_
    ↪already
    insurance_train_data = pd.get_dummies(insurance_train_data,
    ↪columns=["Customer Feedback"], drop_first=False)
    print("Customer Feedback column has been one-hot encoded.")
else:
    print("Customer Feedback column has already been one-hot encoded.")

# Convert the boolean True/False values to 1/0
if 'Customer Feedback_Poor' in insurance_train_data.columns:
    insurance_train_data['Customer Feedback_Poor'] =
    ↪insurance_train_data['Customer Feedback_Poor'].astype(int)
if 'Customer Feedback_Average' in insurance_train_data.columns:
    insurance_train_data['Customer Feedback_Average'] =
    ↪insurance_train_data['Customer Feedback_Average'].astype(int)
if 'Customer Feedback_Good' in insurance_train_data.columns:
    insurance_train_data['Customer Feedback_Good'] =
    ↪insurance_train_data['Customer Feedback_Good'].astype(int)
if 'Customer Feedback_Not Provided' in insurance_train_data.columns:

```



```
insurance_train_data['Customer Feedback_Not Provided'] =
↳insurance_train_data['Customer Feedback_Not Provided'].astype(int)
```

Customer Feedback column has been one-hot encoded.

```
[56]: # 18. Smoking Status - No = 0; Yes = 1
# Map "No" to 0 and "Yes" to 1 in the Smoking Status column
insurance_train_data['Smoking Status'] = insurance_train_data['Smoking Status'].
↳map({'No': 0, 'Yes': 1})
# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Fit and transform the 'Smoking Status' column to convert "No" and "Yes" to 0
↳and 1
insurance_train_data['Smoking Status'] = label_encoder.
↳fit_transform(insurance_train_data['Smoking Status'])
# Check the result
print(insurance_train_data['Smoking Status'].value_counts())
```

Smoking Status

1 570424

0 566662

Name: count, dtype: int64

```
[57]: # 19. Exercise Frequency - Label Encoding based on frequency
# Initialize the LabelEncoder
label_encoder = LabelEncoder()

# Define the desired order for the 'Exercise Frequency' categories
desired_order = ["Not Provided", "Rarely", "Daily", "Weekly", "Monthly"]

# Ensure the encoder fits based on the custom order
label_encoder.classes_ = np.array(desired_order)

# Apply the label encoder
insurance_train_data['Exercise Frequency Encoded'] = label_encoder.
↳transform(insurance_train_data['Exercise Frequency'])
# Check the result
print(insurance_train_data['Exercise Frequency Encoded'].value_counts())
```

Exercise Frequency Encoded

3 290106

4 284292

1 283880

2 278808

Name: count, dtype: int64

```
[58]: # 20. Property Type - One-Hot Encoding
# Check if the "Property Type" column still exists in the DataFrame
if "Property Type" in insurance_train_data.columns:
    # Perform one-hot encoding for "Customer Feedback" if it hasn't been done
    ↪ already
    insurance_train_data = pd.get_dummies(insurance_train_data,
    ↪ columns=["Property Type"], drop_first=False)
    print("Property Type column has been one-hot encoded.")
else:
    print("Property Type column has already been one-hot encoded.")

# Convert the boolean True/False values to 1/0
if 'Property Type_House' in insurance_train_data.columns:
    insurance_train_data['Property Type_House'] =
    ↪ insurance_train_data['Property Type_House'].astype(int)
if 'Property Type_Apartment' in insurance_train_data.columns:
    insurance_train_data['Property Type_Apartment'] =
    ↪ insurance_train_data['Property Type_Apartment'].astype(int)
if 'Property Type_Condo' in insurance_train_data.columns:
    insurance_train_data['Property Type_Condo'] =
    ↪ insurance_train_data['Property Type_Condo'].astype(int)
if 'Property Type_Unknown' in insurance_train_data.columns:
    insurance_train_data['Property Type_Unknown'] =
    ↪ insurance_train_data['Property Type_Unknown'].astype(int)
```

Property Type column has been one-hot encoded.

```
[59]: # Check ideal Lambda value for Premium Amount
premium_amount_data = insurance_train_data['Premium Amount']

# Apply Box-Cox transformation and calculate the optimal lambda
transformed_data, optimal_lambda = boxcox(premium_amount_data)

# Plot the transformed data for a range of lambdas
lambda_range = np.linspace(-2, 2, 100) # Lambda range from -2 to 2
transformed_range = np.array([boxcox(premium_amount_data, lmbda=lmbda)[0] for
    ↪ lmbda in lambda_range])

# Output optimal lambda and the range used for transformations
print(f"Optimal lambda: {optimal_lambda}")
print(f"Lambda range: {lambda_range[0]} to {lambda_range[-1]}")
```

Optimal lambda: 0.4061359533323342

Lambda range: -2.0 to 2.0

```
[60]: # 21. Premium Amount - Apply Box-Cox Transformation (to be reverted after
    ↪ modelling)
# Apply Box-Cox transformation with a hardcoded lambda value
```

```

lambda_value = 0.406
insurance_train_data['Premium Amount'] = boxcox(insurance_train_data['Premium_
↪Amount'] + 1, lmbda=lambda_value)

# Check the result
print(insurance_train_data['Premium Amount'].value_counts())
print(f"Lambda value for 'Premium Amount': {lambda_value}")

```

```

Premium Amount
6.782930      4022
6.636867      3662
6.014977      3628
6.487290      3308
7.202073      3240
...
75.694947         1
73.805532         1
75.350581         1
72.524887         1
75.273752         1
Name: count, Length: 4787, dtype: int64
Lambda value for 'Premium Amount': 0.406

```

Handling Outliers

```

[61]: # Numerical features that will not require outlier handling:
# 1. Age underwent binning
# 2. Annual Income underwent binning

```

```

[62]: # Number of Dependents - Cap to a maximum of 10
insurance_train_data['Number of Dependents'] = np.where(
    insurance_train_data['Number of Dependents'] > 10, 10,
    insurance_train_data['Number of Dependents']
)
print(insurance_train_data['Number of Dependents'].describe())

```

```

count      1.137086e+06
mean       1.826128e+00
std        1.469595e+00
min        0.000000e+00
25%        0.000000e+00
50%        2.000000e+00
75%        3.000000e+00
max        4.000000e+00
Name: Number of Dependents, dtype: float64

```

```

[63]: # Health Score - Drop extreme values
# Calculate the first and third quartile (Q1 and Q3)
Q1 = insurance_train_data['Health Score'].quantile(0.25)

```

```

Q3 = insurance_train_data['Health Score'].quantile(0.75)

# Calculate the Interquartile Range (IQR)
IQR = Q3 - Q1

# Define the lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Filter out the rows where 'Health Score' is an outlier
insurance_train_data = insurance_train_data[(insurance_train_data['Health_
Score'] >= lower_bound) &
                                             (insurance_train_data['Health_
Score'] <= upper_bound)]
print(insurance_train_data['Health Score'].describe())

```

```

count      1.137086e+06
mean       2.398984e+01
std        1.330780e+01
min        0.000000e+00
25%        1.401079e+01
50%        2.326426e+01
75%        3.373437e+01
max        5.897591e+01
Name: Health Score, dtype: float64

```

```

[64]: # Previous Claims - Undergo binning
# Define the bins and labels for Previous Claims based on the stats
bins = [-1, 0, 2, 5, 9, float('inf')] # Define the bin edges
labels = ['No Claims', 'Few Claims (1-2)', 'Moderate Claims (3-5)', 'High_
Claims (6-9)', 'Extreme Claims (9+)'] # Define the bin labels
# Apply binning using pd.cut
insurance_train_data['Previous Claims_Bin'] = pd.
    cut(insurance_train_data['Previous Claims'], bins=bins, labels=labels)
# Check the result
print(insurance_train_data['Previous Claims_Bin'].value_counts())

# Define the mapping of age bin labels to numeric values
previous_claims_bin_mapping = {
    "No Claims": 1,
    "Few Claims (1-2)": 2,
    "Moderate Claims (3-5)": 3,
    "High Claims (6-9)": 4,
    "Extreme Claims (9+)": 5
}
# Apply the mapping to the Previous Claims Bin column to convert it to numeric_
values

```

```

insurance_train_data["Previous Claims_Bin Numeric"] =
    ↪insurance_train_data["Previous Claims_Bin"].map(previous_claims_bin_mapping)
# Convert the binned column to integer codes
insurance_train_data["Previous Claims_Bin Numeric"] =
    ↪insurance_train_data["Previous Claims_Bin Numeric"].cat.codes
# Check the result
print(insurance_train_data[['Previous Claims', 'Previous Claims_Bin', 'Previous_
    ↪Claims_Bin Numeric']].head())

```

```

Previous Claims_Bin
No Claims          635423
Few Claims (1-2)    501663
Moderate Claims (3-5)  0
High Claims (6-9)    0
Extreme Claims (9+)  0
Name: count, dtype: int64

   Previous Claims  Previous Claims_Bin  Previous Claims_Bin Numeric
0              0.605892      Few Claims (1-2)                      1
1              0.468220      Few Claims (1-2)                      1
2              0.468220      Few Claims (1-2)                      1
3              0.468220      Few Claims (1-2)                      1
4              0.000000           No Claims                        0

```

```

[65]: # Vehicle Age - Winsorization/Cap to a maximum of 15
print("Before:")
print(insurance_train_data['Vehicle Age'].describe())
insurance_train_data['Vehicle Age'] = np.where(
    insurance_train_data['Vehicle Age'] > 15, 15,
    insurance_train_data['Vehicle Age']
)
print("After:")
print(insurance_train_data['Vehicle Age'].describe())

```

```

Before:
count      1.137086e+06
mean       9.571087e+00
std        5.776555e+00
min        0.000000e+00
25%        5.000000e+00
50%        1.000000e+01
75%        1.500000e+01
max        1.900000e+01
Name: Vehicle Age, dtype: float64
After:
count      1.137086e+06
mean       9.061272e+00
std        5.103658e+00
min        0.000000e+00

```

```

25%      5.000000e+00
50%      1.000000e+01
75%      1.500000e+01
max       1.500000e+01
Name: Vehicle Age, dtype: float64

```

```

[66]: # Credit Score - Drop extreme values
# Calculate the first and third quartile (Q1 and Q3)
Q1 = insurance_train_data['Credit Score'].quantile(0.25)
Q3 = insurance_train_data['Credit Score'].quantile(0.75)

# Calculate the Interquartile Range (IQR)
IQR = Q3 - Q1

# Define the lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Filter out the rows where 'Credit Score' is an outlier
insurance_train_data = insurance_train_data[(insurance_train_data['Credit_
↪Score'] >= lower_bound) &
                                                    (insurance_train_data['Credit_
↪Score'] <= upper_bound)]
print(insurance_train_data['Credit Score'].describe())

```

```

count      1.004659e+06
mean       5.961158e+02
std        1.488250e+02
min        3.000000e+02
25%        4.730000e+02
50%        6.000000e+02
75%        7.230000e+02
max        8.490000e+02
Name: Credit Score, dtype: float64

```

```

[67]: # Insurance Duration - Undergo binning for data flexibility
# Define the bins and labels for Insurance Duration
bins = [-1, 1, 3, 8, float('inf')] # Define the bin edges
labels = ['New Clients (0-1)', 'Repeat Clients (2-3)', 'Established Clients_
↪(4-8)', 'Very Loyal Clients (9+)'] # Define the bin labels
# Apply binning using pd.cut
insurance_train_data['Insurance Duration_Bin'] = pd.
↪cut(insurance_train_data['Insurance Duration'], bins=bins, labels=labels)
# Check the result
print(insurance_train_data['Insurance Duration_Bin'].value_counts())

# Define the mapping of age bin labels to numeric values
insurance_duration_bin_mapping = {

```

```

    "New Clients (0-1)": 1,
    "Repeat Clients (2-3)": 2,
    "Established Clients (4-8)": 3,
    "Very Loyal Clients (9+)": 4
}
# Apply the mapping to the Insurance Duration Bin column to convert it to
↳ numeric values
insurance_train_data["Insurance Duration_Bin Numeric"] =
↳ insurance_train_data["Insurance Duration_Bin"].
↳ map(insurance_duration_bin_mapping)
# Convert the binned column to integer codes
insurance_train_data["Insurance Duration_Bin Numeric"] =
↳ insurance_train_data["Insurance Duration_Bin Numeric"].cat.codes
# Check the result
print(insurance_train_data[['Insurance Duration', 'Insurance Duration_Bin',
↳ 'Insurance Duration_Bin Numeric']].head())

```

```

Insurance Duration_Bin
Established Clients (4-8)    556463
Repeat Clients (2-3)        220049
Very Loyal Clients (9+)     115411
New Clients (0-1)           112736
Name: count, dtype: int64

   Insurance Duration  Insurance Duration_Bin \
0                5.0  Established Clients (4-8)
1                2.0    Repeat Clients (2-3)
3                1.0    New Clients (0-1)
4                4.0  Established Clients (4-8)
5                5.0  Established Clients (4-8)

   Insurance Duration_Bin Numeric
0                        2
1                        1
3                        0
4                        2
5                        2

```

```

[68]: # Premium Amount - Drop extreme values
# Calculate the first and third quartile (Q1 and Q3)
Q1 = insurance_train_data['Premium Amount'].quantile(0.25)
Q3 = insurance_train_data['Premium Amount'].quantile(0.75)

# Calculate the Interquartile Range (IQR)
IQR = Q3 - Q1

# Define the lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR

```

```
upper_bound = Q3 + 1.5 * IQR

# Filter out the rows where 'Premium Amount' is an outlier
insurance_train_data = insurance_train_data[(insurance_train_data['Premium_Amount'] >= lower_bound) &
                                             (insurance_train_data['Premium_Amount'] <= upper_bound)]
print(insurance_train_data['Premium Amount'].describe())
```

```
count    1.000534e+06
mean     3.653932e+01
std      1.366349e+01
min      6.014977e+00
25%      2.890800e+01
50%      3.620339e+01
75%      4.550665e+01
max      7.084974e+01
Name: Premium Amount, dtype: float64
```

0.4 Statistical Analysis & Tests

Checking Data Distributions

```
[69]: # Check sample values
insurance_train_data.head()
```

```
[69]:
```

	id	Age	Annual Income	Number of Dependents	Education Level	\
0	0	19.0	10049.0	1.0	Bachelor's	
1	1	39.0	31678.0	3.0	Master's	
3	3	21.0	141855.0	2.0	Bachelor's	
4	4	21.0	39651.0	1.0	Bachelor's	
5	5	29.0	45963.0	1.0	Bachelor's	

	Health Score	Previous Claims	Vehicle Age	Credit Score	\
0	22.598761	0.605892	15.0	372.0	
1	15.569731	0.468220	12.0	694.0	
3	10.938144	0.468220	0.0	367.0	
4	20.376094	0.000000	8.0	598.0	
5	33.053198	0.605892	4.0	614.0	

	Insurance Duration	...	Customer Feedback_Not Provided	\
0	5.0	...	0	
1	2.0	...	0	
3	1.0	...	0	
4	4.0	...	0	
5	5.0	...	0	

	Customer Feedback_Poor	Exercise Frequency	Encoded Property Type_Apartment	\
0	1	3	0	

1	0	4	0
3	1	2	1
4	1	3	0
5	0	3	0

	Property Type_Condo	Property Type_House	Previous Claims_Bin \
0	0	1	Few Claims (1-2)
1	0	1	Few Claims (1-2)
3	0	0	Few Claims (1-2)
4	0	1	No Claims
5	0	1	Few Claims (1-2)

	Previous Claims_Bin Numeric	Insurance Duration_Bin \
0	1	Established Clients (4-8)
1	1	Repeat Clients (2-3)
3	1	New Clients (0-1)
4	0	Established Clients (4-8)
5	1	Established Clients (4-8)

	Insurance Duration_Bin Numeric
0	2
1	1
3	0
4	2
5	2

[5 rows x 46 columns]

```
[70]: insurance_train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1000534 entries, 0 to 1199996
Data columns (total 46 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    1000534 non-null  int64
1   Age                                  1000534 non-null  float64
2   Annual Income                        1000534 non-null  float64
3   Number of Dependents                 1000534 non-null  float64
4   Education Level                      1000534 non-null  object
5   Health Score                         1000534 non-null  float64
6   Previous Claims                     1000534 non-null  float64
7   Vehicle Age                         1000534 non-null  float64
8   Credit Score                        1000534 non-null  float64
9   Insurance Duration                  1000534 non-null  float64
10  Smoking Status                      1000534 non-null  int64
11  Exercise Frequency                  1000534 non-null  object
12  Premium Amount                     1000534 non-null  float64
```

13	Age_Bin	1000534	non-null	category
14	Age_Bin Numeric	1000534	non-null	int8
15	Gender_Female	1000534	non-null	int64
16	Gender_Male	1000534	non-null	int64
17	Income_Bin	1000534	non-null	category
18	Income_Bin Numeric	1000534	non-null	int8
19	Marital Status_Divorced	1000534	non-null	bool
20	Marital Status_Married	1000534	non-null	bool
21	Marital Status_Single	1000534	non-null	bool
22	Education Level_Encoded	1000534	non-null	int64
23	Occupation_Employed/Self-Employed	1000534	non-null	int64
24	Occupation_Unemployed	1000534	non-null	int64
25	Location_Rural	1000534	non-null	int64
26	Location_Suburban	1000534	non-null	int64
27	Location_Urban	1000534	non-null	int64
28	Policy Type_Basic	1000534	non-null	int64
29	Policy Type_Comprehensive	1000534	non-null	int64
30	Policy Type_Premium	1000534	non-null	int64
31	Policy Start Year	1000534	non-null	int32
32	Policy Start Month	1000534	non-null	int32
33	Years Since Start	1000534	non-null	float64
34	Customer Feedback_Average	1000534	non-null	int64
35	Customer Feedback_Good	1000534	non-null	int64
36	Customer Feedback_Not Provided	1000534	non-null	int64
37	Customer Feedback_Poor	1000534	non-null	int64
38	Exercise Frequency Encoded	1000534	non-null	int64
39	Property Type_Apartment	1000534	non-null	int64
40	Property Type_Condo	1000534	non-null	int64
41	Property Type_House	1000534	non-null	int64
42	Previous Claims_Bin	1000534	non-null	category
43	Previous Claims_Bin Numeric	1000534	non-null	int8
44	Insurance Duration_Bin	1000534	non-null	category
45	Insurance Duration_Bin Numeric	1000534	non-null	int8

dtypes: bool(3), category(4), float64(10), int32(2), int64(21), int8(4), object(2)

memory usage: 277.7+ MB

```
[71]: # Histogram (For Overall Distribution)
#insurance_train_data.hist(figsize=(12, 10), bins=30, edgecolor='black')
#plt.suptitle("Updated Feature Distributions", fontsize=16)
#plt.show()

# Check distribution of numerical features
# Get numerical columns
num_columns = insurance_train_data.select_dtypes(include=['number']).columns

# Define grid size (5 columns)
```

```

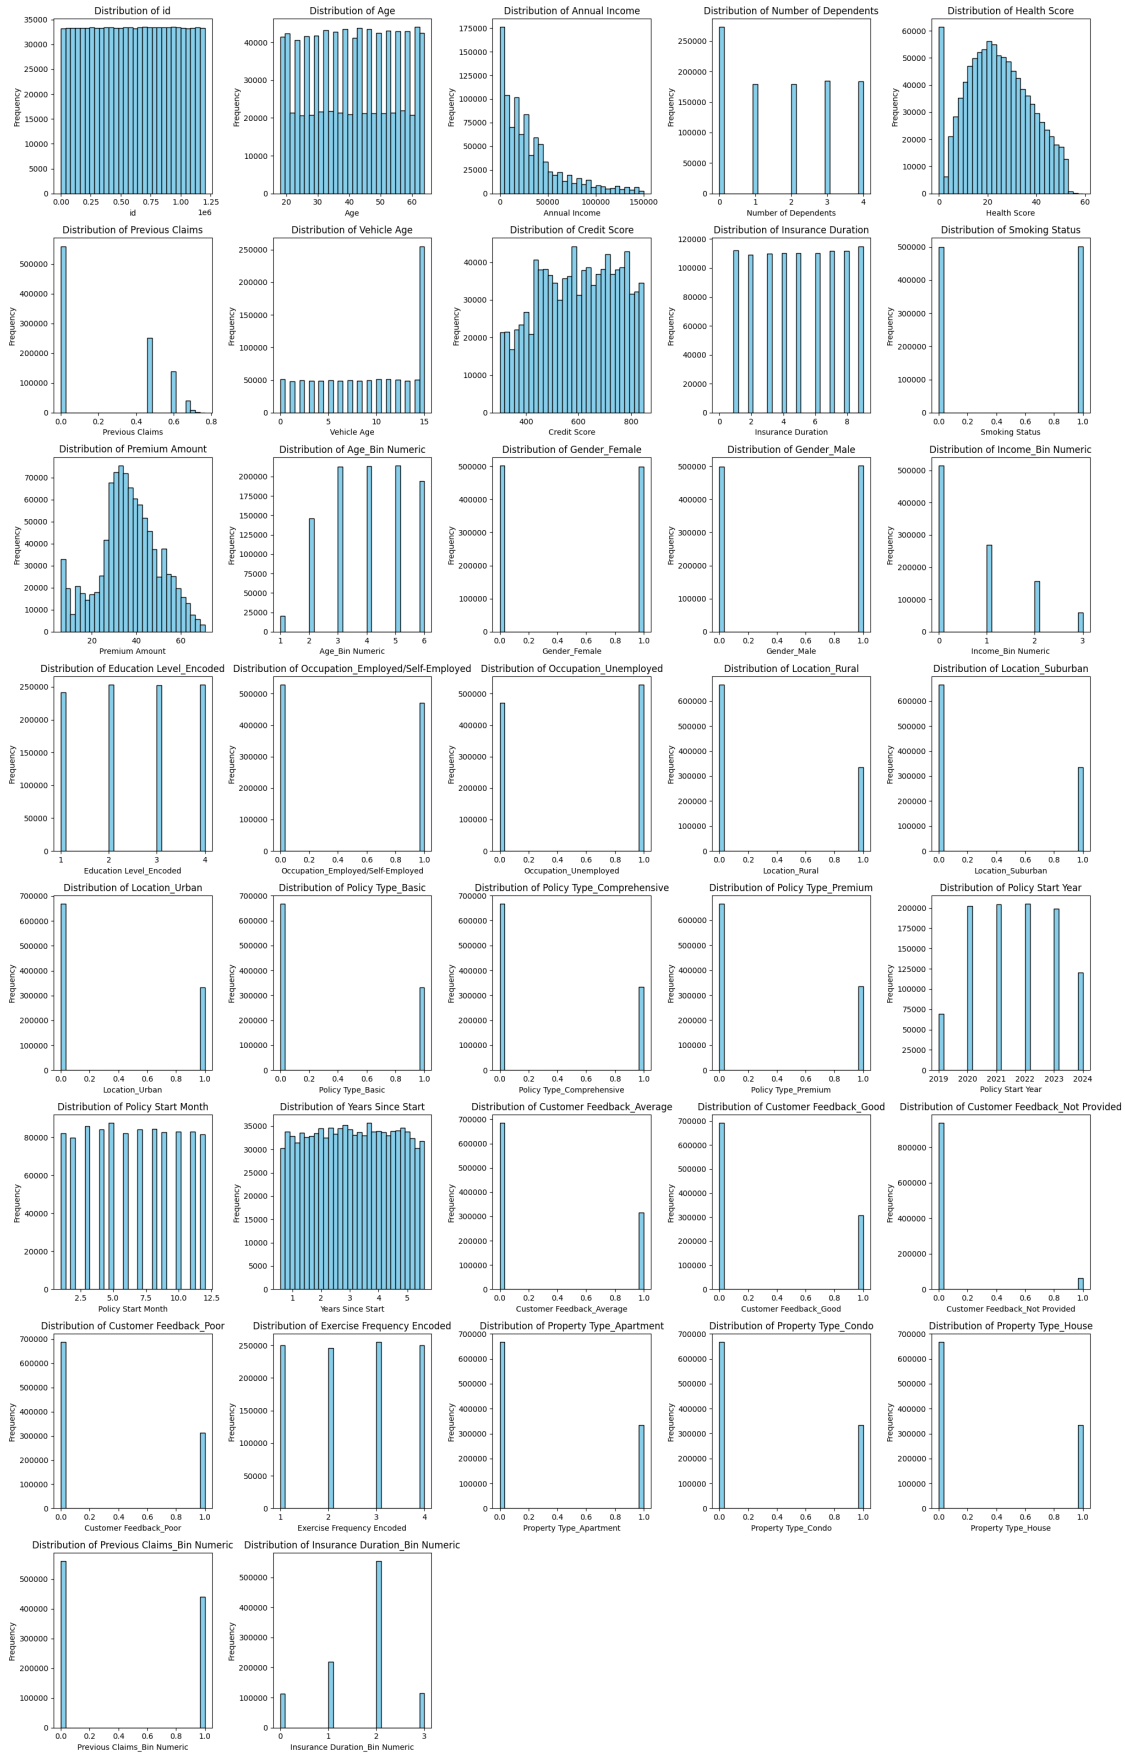
num_cols = 5
num_rows = (len(num_columns) // num_cols) + (len(num_columns) % num_cols > 0)

# Create a figure with a grid of subplots
plt.figure(figsize=(20, 4 * num_rows)) # Adjust height based on the number of
    ↪ rows

# Loop through numerical columns and plot histograms
for i, col in enumerate(num_columns, 1):
    plt.subplot(num_rows, num_cols, i) # Position plot in a 4x5 grid
    plt.hist(insurance_train_data[col], bins=30, color='skyblue',
    ↪ edgecolor='black')
    plt.title(f"Distribution of {col}")
    plt.xlabel(col)
    plt.ylabel("Frequency")

# Adjust layout to prevent overlap
plt.tight_layout()
plt.show()

```



```
[72]: # Skewness Computation with Interpretation
# Select only numerical columns
numerical_columns = insurance_train_data.select_dtypes(include=['number'])

# Compute skewness
skew_values = numerical_columns.skew().sort_values(ascending=False)

# Create a DataFrame to store skewness values
skew_df = pd.DataFrame({'Skewness Value': skew_values})

# Apply interpretation directly
skew_df['Interpretation'] = ""

skew_df.loc[skew_df['Skewness Value'] == 0, 'Interpretation'] = "Norm Dist ( = 0)"
skew_df.loc[(skew_df['Skewness Value'] > -0.5) & (skew_df['Skewness Value'] < 0.5), 'Interpretation'] = "Min/No Skew ( = -0.5 to 0.5)"
skew_df.loc[skew_df['Skewness Value'] > 1, 'Interpretation'] = "Extreme Right-Skewed ( > 1)"
skew_df.loc[skew_df['Skewness Value'] < -1, 'Interpretation'] = "Extreme Left-Skewed ( < -1)"
skew_df.loc[(skew_df['Skewness Value'] >= 0.5) & (skew_df['Skewness Value'] <= 1), 'Interpretation'] = "Right-Skewed ( > 0.5)"
skew_df.loc[(skew_df['Skewness Value'] <= -0.5) & (skew_df['Skewness Value'] >= -1), 'Interpretation'] = "Left-Skewed ( < -0.5)"

# Display the skewness table
print(skew_df)
```

	Skewness Value \
Customer Feedback_Not Provided	3.584894
Annual Income	1.429848
Income_Bin Numeric	0.945397
Customer Feedback_Good	0.830539
Customer Feedback_Poor	0.807069
Customer Feedback_Average	0.793994
Location_Urban	0.717122
Policy Type_Basic	0.711230
Policy Type_Comprehensive	0.708570
Property Type_Apartment	0.707582
Property Type_House	0.707343
Property Type_Condo	0.706399
Location_Suburban	0.702387
Location_Rural	0.701849
Policy Type_Premium	0.701535

Previous Claims	0.362348
Previous Claims_Bin Numeric	0.240599
Number of Dependents	0.132191
Health Score	0.119646
Occupation_Employed/Self-Employed	0.115812
Gender_Female	0.008643
Policy Start Month	0.006429
id	-0.000634
Smoking Status	-0.005341
Policy Start Year	-0.005895
Years Since Start	-0.006044
Gender_Male	-0.008643
Insurance Duration	-0.010852
Age	-0.011916
Exercise Frequency Encoded	-0.012130
Education Level_Encoded	-0.018945
Premium Amount	-0.100817
Occupation_Unemployed	-0.115812
Credit Score	-0.134270
Age_Bin Numeric	-0.139107
Vehicle Age	-0.294088
Insurance Duration_Bin Numeric	-0.543191

	Interpretation
Customer Feedback_Not Provided	Extreme Right-Skewed (> 1)
Annual Income	Extreme Right-Skewed (> 1)
Income_Bin Numeric	Right-Skewed (> 0.5)
Customer Feedback_Good	Right-Skewed (> 0.5)
Customer Feedback_Poor	Right-Skewed (> 0.5)
Customer Feedback_Average	Right-Skewed (> 0.5)
Location_Urban	Right-Skewed (> 0.5)
Policy Type_Basic	Right-Skewed (> 0.5)
Policy Type_Comprehensive	Right-Skewed (> 0.5)
Property Type_Apartment	Right-Skewed (> 0.5)
Property Type_House	Right-Skewed (> 0.5)
Property Type_Condo	Right-Skewed (> 0.5)
Location_Suburban	Right-Skewed (> 0.5)
Location_Rural	Right-Skewed (> 0.5)
Policy Type_Premium	Right-Skewed (> 0.5)
Previous Claims	Min/No Skew (= -0.5 to 0.5)
Previous Claims_Bin Numeric	Min/No Skew (= -0.5 to 0.5)
Number of Dependents	Min/No Skew (= -0.5 to 0.5)
Health Score	Min/No Skew (= -0.5 to 0.5)
Occupation_Employed/Self-Employed	Min/No Skew (= -0.5 to 0.5)
Gender_Female	Min/No Skew (= -0.5 to 0.5)
Policy Start Month	Min/No Skew (= -0.5 to 0.5)
id	Min/No Skew (= -0.5 to 0.5)
Smoking Status	Min/No Skew (= -0.5 to 0.5)

Policy Start Year	Min/No Skew (= -0.5 to 0.5)
Years Since Start	Min/No Skew (= -0.5 to 0.5)
Gender_Male	Min/No Skew (= -0.5 to 0.5)
Insurance Duration	Min/No Skew (= -0.5 to 0.5)
Age	Min/No Skew (= -0.5 to 0.5)
Exercise Frequency Encoded	Min/No Skew (= -0.5 to 0.5)
Education Level_Encoded	Min/No Skew (= -0.5 to 0.5)
Premium Amount	Min/No Skew (= -0.5 to 0.5)
Occupation_Unemployed	Min/No Skew (= -0.5 to 0.5)
Credit Score	Min/No Skew (= -0.5 to 0.5)
Age_Bin Numeric	Min/No Skew (= -0.5 to 0.5)
Vehicle Age	Min/No Skew (= -0.5 to 0.5)
Insurance Duration_Bin Numeric	Left-Skewed (< -0.5)

```
[73]: # Boxplots for outlier detection

# Get numerical columns
num_columns = insurance_train_data.select_dtypes(include=['number']).columns

# Define grid size (5 columns)
num_cols = 5
num_rows = (len(num_columns) // num_cols) + (len(num_columns) % num_cols > 0)

# Create a figure with a grid of subplots
plt.figure(figsize=(20, 4 * num_rows)) # Adjust height based on the number of
↪ rows

# Loop through numerical columns and create boxplots
for i, col in enumerate(num_columns, 1):
    plt.subplot(num_rows, num_cols, i) # Position plot in a 4x5 grid
    sns.boxplot(y=insurance_train_data[col])
    plt.title(f"Boxplot of {col}")

# Adjust layout to prevent overlap
plt.tight_layout()
plt.show()
```



Feature Selection

```
[74]: # Multicollinearity Check (VIF Analysis)
# Step 1: Prepare dataset (select numerical features only)
X = insurance_train_data.select_dtypes(include=['number']).
    ↪ drop(columns=['Premium Amount'])

# Step 2: Exclude one-hot encoded columns (e.g., columns related to "Gender" or
    ↪ other categorical variables)
one_hot_columns = [col for col in X.columns if 'Gender' in col or 'Marital_
    ↪ Status' in col
                    or 'Occupation' in col or 'Location' in col or 'Policy_
    ↪ Type' in col
                    or 'Customer Feedback' in col or 'Property Type' in col]
    ↪ # Add more as needed
binned_columns = [col for col in X.columns if 'Bin' in col or 'Numeric' in col]
    ↪ # Include binned features
time_related_columns = ['Policy Start Year', 'Years Since Start'] #
    ↪ Time-related features that may cause collinearity

# Combine all features to exclude
exclude_columns = one_hot_columns + binned_columns + time_related_columns

# Drop one-hot encoded columns
X = X.drop(columns=exclude_columns)

# Step 3: Add a constant column to the dataset (for intercept)
X_const = add_constant(X)

# Step 4: Calculate VIF for each feature
vif_data = pd.DataFrame()
vif_data["Variable"] = X_const.columns
vif_data["VIF"] = [variance_inflation_factor(X_const.values, i) for i in
    ↪ range(X_const.shape[1])]

# Step 5: Display VIF results
print("VIF Analysis:\n", vif_data)

# Step 6: Identify features with high VIF (typically > 5 or 10)
high_vif = vif_data[vif_data['VIF'] > 5] # or 10 depending on your threshold
print("\nFeatures with high VIF:\n", high_vif)
```

VIF Analysis:

	Variable	VIF
0	const	59.478027

1	id	1.000007
2	Age	1.000026
3	Annual Income	1.044218
4	Number of Dependents	1.000043
5	Health Score	1.000712
6	Previous Claims	1.001988
7	Vehicle Age	1.000028
8	Credit Score	1.044025
9	Insurance Duration	1.000022
10	Smoking Status	1.000020
11	Education Level_Encoded	1.000020
12	Policy Start Month	1.000115
13	Exercise Frequency Encoded	1.000027

Features with high VIF:

	Variable	VIF
0	const	59.478027

```
[75]: # Correlation Analysis (v2)
# Select only numerical columns
num_features = insurance_train_data.select_dtypes(include=['number'])

# Compute correlation matrix
corr_matrix = num_features.corr()

# Identify highly correlated features (threshold > 0.9 or a threshold of your
↳choice)
threshold = 0.9
highly_corr_features = set()

for i in range(len(corr_matrix.columns)):
    for j in range(i):
        if abs(corr_matrix.iloc[i, j]) > threshold:
            colname = corr_matrix.columns[i]
            highly_corr_features.add(colname)

print(f"Highly correlated features (threshold {threshold}): \n",
↳highly_corr_features)
```

Highly correlated features (threshold 0.9):

```
{'Income_Bin Numeric', 'Age_Bin Numeric', 'Previous Claims_Bin Numeric',
'Occupation_Unemployed', 'Gender_Male', 'Years Since Start'}
```

```
[76]: # Feature Importance Using Random Forest
# Select only numerical columns
num_features = insurance_train_data.select_dtypes(include=['number'])

# Prepare the dataset (exclude target column and select only numerical features)
```

```

X = num_features.select_dtypes(include=['number']).drop(columns=['Premium_Amount']) # Select only numerical columns
y = num_features['Premium Amount'] # Target variable

# Train a Random Forest model
# rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf = RandomForestRegressor(n_estimators=50, random_state=42, n_jobs=-1,
    max_features='sqrt')
rf.fit(X, y)

# Get feature importance scores
feature_importance = pd.Series(rf.feature_importances_, index=X.columns)

# Sort the features by importance
top_features = feature_importance.sort_values(ascending=False)

# Display top features
print("Top Features by Importance:\n", top_features.head(10))

```

Top Features by Importance:

Credit Score	0.103112
Annual Income	0.097741
Health Score	0.092676
id	0.083900
Years Since Start	0.081385
Age	0.062483
Vehicle Age	0.048043
Policy Start Month	0.044102
Insurance Duration	0.037889
Number of Dependents	0.029792

dtype: float64

```

[77]: # Recursive Feature Elimination (RFE)
# Select only numerical columns
num_features = insurance_train_data.select_dtypes(include=['number'])

# Step 1: Prepare the dataset
X = num_features.drop(columns=['Premium Amount']) # Features
y = num_features['Premium Amount'] # Target variable

# Step 3: Identify categorical columns for encoding
categorical_columns = X.select_dtypes(include=['object']).columns

# Create a column transformer to apply one-hot encoding to categorical features
preprocessor = ColumnTransformer(
    transformers=[('cat', OneHotEncoder(drop='first'), categorical_columns)],
    remainder='passthrough' # Keep numerical columns as they are

```

```

)

# Step 3: Apply RFE for feature selection
# Create a pipeline with preprocessing and feature selection
pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor), # Preprocessing step
    ('scaler', StandardScaler()), # Optional: Scale features before RFE
    ('feature_selection', RFE(LinearRegression(), n_features_to_select=10)) #_
    ↪ RFE feature selection
])

# Fit the pipeline
pipeline.fit(X, y)

# Step 4: Get selected features
# Get the selected features after RFE
selected_features = X.columns[pipeline.named_steps['feature_selection'].
    ↪ support_]

print("Selected Features:\n", selected_features)

```

Selected Features:

```

Index(['Annual Income', 'Previous Claims', 'Credit Score',
      'Income_Bin Numeric', 'Occupation_Unemployed', 'Policy Start Year',
      'Policy Start Month', 'Years Since Start',
      'Customer Feedback_Not Provided', 'Previous Claims_Bin Numeric'],
      dtype='object')

```

```

[78]: # Extract feature importance data
ranking = pipeline.named_steps['feature_selection'].ranking_
support = pipeline.named_steps['feature_selection'].support_

# Create DataFrame for analysis
feature_df = pd.DataFrame({
    'Feature': X.columns,
    'Selected': support,
    'Ranking': ranking
})

# Sort by ranking (1 = most important)
feature_df = feature_df.sort_values(by='Ranking')

# Display the full matrix
print("Feature Ranking Table:")
display(feature_df) # Use in Jupyter Notebook for better output

# Highlight selected features

```

```
print("\nTop Selected Features:")
display(feature_df[feature_df['Selected'] == True])
```

Feature Ranking Table:

	Feature	Selected	Ranking
2	Annual Income	True	1
5	Previous Claims	True	1
7	Credit Score	True	1
13	Income_Bin Numeric	True	1
24	Policy Start Month	True	1
25	Years Since Start	True	1
23	Policy Start Year	True	1
16	Occupation_Unemployed	True	1
28	Customer Feedback_Not Provided	True	1
34	Previous Claims_Bin Numeric	True	1
10	Age_Bin Numeric	False	2
1	Age	False	3
26	Customer Feedback_Average	False	4
27	Customer Feedback_Good	False	5
29	Customer Feedback_Poor	False	6
3	Number of Dependents	False	7
15	Occupation_Employed/Self-Employed	False	8
8	Insurance Duration	False	9
35	Insurance Duration_Bin Numeric	False	10
17	Location_Rural	False	11
30	Exercise Frequency Encoded	False	12
14	Education Level_Encoded	False	13
9	Smoking Status	False	14
22	Policy Type_Premium	False	15
32	Property Type_Condo	False	16
6	Vehicle Age	False	17
18	Location_Suburban	False	18
19	Location_Urban	False	19
4	Health Score	False	20
33	Property Type_House	False	21
20	Policy Type_Basic	False	22
21	Policy Type_Comprehensive	False	23
12	Gender_Male	False	24
31	Property Type_Apartment	False	25
11	Gender_Female	False	26
0	id	False	27

Top Selected Features:

	Feature	Selected	Ranking
2	Annual Income	True	1
5	Previous Claims	True	1

7	Credit Score	True	1
13	Income_Bin Numeric	True	1
24	Policy Start Month	True	1
25	Years Since Start	True	1
23	Policy Start Year	True	1
16	Occupation_Unemployed	True	1
28	Customer Feedback_Not Provided	True	1
34	Previous Claims_Bin Numeric	True	1

0.5 Model Development & Evaluation

Split the Data into Training and Test Sets

```
[79]: # Split the Data into Training and Test Sets (v2)
# Split the data into features (X) and target variable (y)
# Selected Features:
# Index(['Age', 'Annual Income', 'Previous Claims', 'Credit Score',
#        'Age_Bin Numeric', 'Income_Bin Numeric', 'Policy Start Year',
#        'Years Since Start', 'Customer Feedback_Not Provided',
#        'Previous Claims_Bin Numeric'],
#        dtype='object')
# The following feature was excluded as it did not make sense: Customer_
↳Feedback_Not Provided
X = insurance_train_data[['Age', 'Annual Income', 'Previous Claims', 'Credit_
↳Score',
                        'Age_Bin Numeric', 'Income_Bin Numeric', 'Policy_
↳Start Year',
                        'Years Since Start', 'Previous Claims_Bin Numeric']]
y = insurance_train_data['Premium Amount']

# Split data into training and test sets (80-20 split)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

Baseline Models

```
[80]: # Basic Modeling via Multiple Linear Regression
# Train the model
linear_reg = LinearRegression()
linear_reg.fit(X_train, y_train)

# Predict on validation set
y_train_pred = linear_reg.predict(X_train)
y_test_pred = linear_reg.predict(X_test)

# Apply log transformation to actual and predicted values for training and test
log_y_train_true = np.log1p(y_train) # log1p is equivalent to log(y + 1)
log_y_train_pred = np.log1p(y_train_pred)
log_y_test_true = np.log1p(y_test)
```

```

log_y_test_pred = np.log1p(y_test_pred)

# Calculate squared differences between the log-transformed actual and
# predicted values
train_squared_diff = (log_y_train_true - log_y_train_pred) ** 2
test_squared_diff = (log_y_test_true - log_y_test_pred) ** 2

# Evaluate the model on both training and test sets
train_mean_squared_diff = np.mean(train_squared_diff)
train_rmsle = np.sqrt(train_mean_squared_diff)
train_mse = mean_squared_error(y_train, y_train_pred)
train_r2 = r2_score(y_train, y_train_pred)
test_mean_squared_diff = np.mean(test_squared_diff)
test_rmsle = np.sqrt(test_mean_squared_diff)
test_mse = mean_squared_error(y_test, y_test_pred)
test_r2 = r2_score(y_test, y_test_pred)

# Print evaluation metrics for both training and test sets
print("Training Set Evaluation:")
print(f"Mean Squared Error: {train_mse}")
print(f"R-squared: {train_r2}")
print(f"RMSLE: {train_rmsle}")

print("\nTest Set Evaluation:")
print(f"Mean Squared Error: {test_mse}")
print(f"R-squared: {test_r2}")
print(f"RMSLE: {test_rmsle}")

# RMSLE Value      Interpretation
# 0.0 - 0.2        Excellent - Predictions are very close to actual values.
# 0.2 - 0.5        Good - Model performs well, but some errors exist.
# 0.5 - 1.0        Moderate - The model makes significant errors. May need
#                  improvement.
# > 1.0           Poor - Large errors; model likely overfitting or missing key
#                  features.

```

Training Set Evaluation:

Mean Squared Error: 184.91404113383106

R-squared: 0.00917630503125566

RMSLE: 0.47140170390376634

Test Set Evaluation:

Mean Squared Error: 185.3016256200385

R-squared: 0.00879362007160911

RMSLE: 0.47250227456368893

Advanced Models

```

[81]: # Decision Tree
# Train the model
dt_model = DecisionTreeRegressor(random_state=42)
dt_model.fit(X_train, y_train)

# Predict on training set
y_pred_train_dt = dt_model.predict(X_train)
# Predict on validation set (testing set)
y_pred_test_dt = dt_model.predict(X_test)

# Apply log transformation to actual and predicted values for training set
log_y_true_train_dt = np.log1p(y_train) # log1p is equivalent to log(y + 1)
log_y_pred_train_dt = np.log1p(y_pred_train_dt)
# Apply log transformation to actual and predicted values for testing set
log_y_true_test_dt = np.log1p(y_test)
log_y_pred_test_dt = np.log1p(y_pred_test_dt)

# Calculate squared differences between the log-transformed actual and
# predicted values (training set)
squared_diff_train_dt = (log_y_true_train_dt - log_y_pred_train_dt) ** 2
# Calculate squared differences between the log-transformed actual and
# predicted values (testing set)
squared_diff_test_dt = (log_y_true_test_dt - log_y_pred_test_dt) ** 2

# Evaluate the model for training set
mean_squared_diff_train_dt = np.mean(squared_diff_train_dt)
rmsle_train_dt = np.sqrt(mean_squared_diff_train_dt)
mse_train_dt = mean_squared_error(y_train, y_pred_train_dt)
r2_train_dt = r2_score(y_train, y_pred_train_dt)
# Evaluate the model for testing set
mean_squared_diff_test_dt = np.mean(squared_diff_test_dt)
rmsle_test_dt = np.sqrt(mean_squared_diff_test_dt)
mse_test_dt = mean_squared_error(y_test, y_pred_test_dt)
r2_test_dt = r2_score(y_test, y_pred_test_dt)

# Print results
print(f"Decision Tree - Training MSE: {mse_train_dt}")
print(f"Decision Tree - Training R-squared: {r2_train_dt}")
print(f"Training RMSLE: {rmsle_train_dt}")

print(f"Decision Tree - Testing MSE: {mse_test_dt}")
print(f"Decision Tree - Testing R-squared: {r2_test_dt}")
print(f"Testing RMSLE: {rmsle_test_dt}")

```

```

Decision Tree - Training MSE: 0.0257776636001459
Decision Tree - Training R-squared: 0.999861875714038
Training RMSLE: 0.005687402361874336
Decision Tree - Testing MSE: 374.4159753604046

```


Decision Tree - Testing R-squared: -1.0028075969788528
Testing RMSLE: 0.6602510840912299

```
[82]: # Random Forest
# Train the model
# rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
# Optimized Random Forest (if needed)
# Train the model
rf_model = RandomForestRegressor(
    n_estimators=50,      # Reduce number of trees (default: 100)
    max_depth=6,         # Limit tree depth (prevents deep, slow trees)
    min_samples_split=10, # Avoid unnecessary small splits
    min_samples_leaf=4,   # Avoid overfitting, improves efficiency
    max_features="sqrt",  # Use fewer features per split (reduces computation)
    bootstrap=True,       # Ensures faster training with sampling
    n_jobs=1,             # Avoids high CPU load (single-threaded)
    random_state=42
)

# Fit the model on the training data
rf_model.fit(X_train, y_train)

# Predict on the training set
y_pred_train_rf = rf_model.predict(X_train)
# Predict on the test set
y_pred_test_rf = rf_model.predict(X_test)

# Apply log transformation to actual and predicted values for training set
log_y_true_train_rf = np.log1p(y_train) # log1p is equivalent to log(y + 1)
log_y_pred_train_rf = np.log1p(y_pred_train_rf)
# Apply log transformation to actual and predicted values for test set
log_y_true_test_rf = np.log1p(y_test)
log_y_pred_test_rf = np.log1p(y_pred_test_rf)

# Calculate squared differences between the log-transformed actual and
# predicted values for training set
squared_diff_train_rf = (log_y_true_train_rf - log_y_pred_train_rf) ** 2
# Calculate squared differences between the log-transformed actual and
# predicted values for test set
squared_diff_test_rf = (log_y_true_test_rf - log_y_pred_test_rf) ** 2

# Evaluate the model on the training set
mean_squared_diff_train_rf = np.mean(squared_diff_train_rf)
rmsle_train_rf = np.sqrt(mean_squared_diff_train_rf)
mse_train_rf = mean_squared_error(y_train, y_pred_train_rf)
r2_train_rf = r2_score(y_train, y_pred_train_rf)
# Evaluate the model on the test set
```

```

mean_squared_diff_test_rf = np.mean(squared_diff_test_rf)
rmsle_test_rf = np.sqrt(mean_squared_diff_test_rf)
mse_test_rf = mean_squared_error(y_test, y_pred_test_rf)
r2_test_rf = r2_score(y_test, y_pred_test_rf)

# Print results
print(f"Random Forest - Training MSE: {mse_train_rf}")
print(f"Random Forest - Training R-squared: {r2_train_rf}")
print(f"Training RMSLE: {rmsle_train_rf}")

print(f"Random Forest - Testing MSE: {mse_test_rf}")
print(f"Random Forest - Testing R-squared: {r2_test_rf}")
print(f"Testing RMSLE: {rmsle_test_rf}")

```

```

Random Forest - Training MSE: 179.55403470428467
Random Forest - Training R-squared: 0.03789679236155907
Training RMSLE: 0.4646315542918868
Random Forest - Testing MSE: 179.98005079215227
Random Forest - Testing R-squared: 0.03725952749696049
Testing RMSLE: 0.4658704646094586

```

```

[83]: # Prepare the data for XGBoost
dtrain = xgb.DMatrix(X_train, label=y_train)
dtest = xgb.DMatrix(X_test, label=y_test)

# Set XGBoost parameters
# params = {
#     'objective': 'reg:squarederror',
#     'eval_metric': 'rmse',
#     'max_depth': 6,
#     'eta': 0.1
# }
# Optimized XGBoost parameters (if needed)
params = {
    'objective': 'reg:squarederror',
    'eval_metric': 'rmse',
    'max_depth': 4,          # Reduce tree depth (less memory usage)
    'eta': 0.2,              # Faster learning rate (reduces rounds needed)
    'subsample': 0.8,        # Use 80% of data per boosting round (less
    ↪ computation)
    'colsample_bytree': 0.8 # Use 80% of features per tree (faster training)
}

# Reduce boosting rounds
num_boost_round = 50 # Reduce from 100 to 50
# num_boost_round = 100

```

```

# Train the XGBoost model
xgb_model = xgb.train(params, dtrain, num_boost_round=num_boost_round)

# Make predictions
y_pred_xgb = xgb_model.predict(dtest)
# Make predictions on the train set (for RMSLE calculation on training data)
y_pred_xgb_train = xgb_model.predict(dtrain)

# Evaluate the model on the test set
mse_xgb_test = mean_squared_error(y_test, y_pred_xgb)
r2_xgb_test = r2_score(y_test, y_pred_xgb)

# Calculate RMLSE for the test set
log_y_true_xgb = np.log1p(y_test) # log1p is log(y + 1)
log_y_pred_xgb = np.log1p(y_pred_xgb)

# Squared differences between log-transformed actual and predicted values (test
↪ set)
squared_diff_xgb = (log_y_true_xgb - log_y_pred_xgb) ** 2

# Mean squared difference and RMSLE for the test set
mean_squared_diff_xgb = np.mean(squared_diff_xgb)
rmsle_xgb_test = np.sqrt(mean_squared_diff_xgb)

# Evaluate the model on the train set
mse_xgb_train = mean_squared_error(y_train, y_pred_xgb_train)
r2_xgb_train = r2_score(y_train, y_pred_xgb_train)

# Calculate RMLSE for the train set
log_y_true_xgb_train = np.log1p(y_train) # log1p is log(y + 1)
log_y_pred_xgb_train = np.log1p(y_pred_xgb_train)

# Squared differences between log-transformed actual and predicted values
↪ (train set)
squared_diff_xgb_train = (log_y_true_xgb_train - log_y_pred_xgb_train) ** 2

# Mean squared difference and RMSLE for the train set
mean_squared_diff_xgb_train = np.mean(squared_diff_xgb_train)
rmsle_xgb_train = np.sqrt(mean_squared_diff_xgb_train)

# Print results
print(f"XGBoost Test MSE: {mse_xgb_test}")
print(f"XGBoost Test R-squared: {r2_xgb_test}")
print(f"XGBoost Test RMSLE: {rmsle_xgb_test}")

print(f"XGBoost Train MSE: {mse_xgb_train}")
print(f"XGBoost Train R-squared: {r2_xgb_train}")

```

```
print(f"XGBoost Train RMSLE: {rmsle_xgb_train}")
```

```
XGBoost Test MSE: 179.71247595680464
XGBoost Test R-squared: 0.03869082570072635
XGBoost Test RMSLE: 0.4654407429199889
XGBoost Train MSE: 179.19570658818637
XGBoost Train R-squared: 0.03981681955812422
XGBoost Train RMSLE: 0.4641039906029737
```

Hyperparameter Tuning

```
[84]: # Define models and hyperparameter grids
      # We modified the settings for standard laptop use
models = {
    "Decision Tree": {
        "model": DecisionTreeRegressor(random_state=42),
        "params": {
            'max_depth': [3, 6],
            'min_samples_split': [5, 10],
            'min_samples_leaf': [2, 4],
            'max_features': ['sqrt'], # ['auto', 'sqrt', 'log2'],
            'criterion': ['squared_error'] # ['squared_error', 'absolute_error']
        }
    },
    "Random Forest": {
        "model": RandomForestRegressor(random_state=42),
        "params": {
            'n_estimators': [50, 100],
            'max_depth': [3, 6],
            'min_samples_split': [5, 10],
            'min_samples_leaf': [2, 4],
            'max_features': ['sqrt'], # ['auto', 'sqrt', 'log2'],
            'bootstrap': [True] # [True, False]
        }
    },
    "XGBoost": {
        "model": xgb.XGBRegressor(objective='reg:squarederror',
        eval_metric='rmse'),
        "params": {
            'max_depth': [3, 6],
            'eta': [0.1, 0.2],
            'n_estimators': [50, 100],
            'subsample': [0.8, 1.0],
            'colsample_bytree': [0.8, 1.0],
            'gamma': [0, 0.1]
        }
    }
}
```

```

# Dictionary to store results
results_rsc = {}

# Perform RandomizedSearchCV and evaluate each model
for name, config in models.items():
    print(f"Training {name}...")

    random_search = RandomizedSearchCV(
        config["model"], param_distributions=config["params"],
        n_iter=min(10, len(ParameterGrid(config["params"]))), # Adjust n_iter_
        ↪automatically
        cv=3, verbose=2, random_state=42, n_jobs=-1
    )
    random_search.fit(X_train, y_train)
    best_model = random_search.best_estimator_

    # Predict and evaluate
    y_pred_rsc = best_model.predict(X_test)
    rmse_rsc = np.sqrt(mean_squared_error(y_test, y_pred_rsc))
    r2_rsc = r2_score(y_test, y_pred_rsc)
    rmsle_rsc = np.sqrt(mean_squared_error(np.log1p(y_test), np.
        ↪log1p(y_pred_rsc)))

    # Store results
    results_rsc[name] = {
        "Best Parameters": random_search.best_params_,
        "RMSE": rmse_rsc,
        "RMSLE": rmsle_rsc,
        "R-squared": r2_rsc
    }

# Print results for all models
for name, metrics in results_rsc.items():
    print(f"\n{name} Results:")
    print(f"Best Parameters: {metrics['Best Parameters']}")
    print(f"RMSE: {metrics['RMSE']}")
    print(f"RMSLE: {metrics['RMSLE']}")
    print(f"R-squared: {metrics['R-squared']}")

```

```

Training Decision Tree...
Fitting 3 folds for each of 8 candidates, totalling 24 fits
Training Random Forest...
Fitting 3 folds for each of 10 candidates, totalling 30 fits
Training XGBoost...
Fitting 3 folds for each of 10 candidates, totalling 30 fits

```

Decision Tree Results:

Best Parameters: {'min_samples_split': 10, 'min_samples_leaf': 2, 'max_features': 'sqrt', 'max_depth': 6, 'criterion': 'squared_error'}
RMSE: 13.467007387307165
RMSLE: 0.4675707029107024
R-squared: 0.02987643038862786

Random Forest Results:

Best Parameters: {'n_estimators': 100, 'min_samples_split': 5, 'min_samples_leaf': 2, 'max_features': 'sqrt', 'max_depth': 6, 'bootstrap': True}
RMSE: 13.417428748126161
RMSLE: 0.4659722004214911
R-squared: 0.03700628039945408

XGBoost Results:

Best Parameters: {'subsample': 0.8, 'n_estimators': 50, 'max_depth': 6, 'gamma': 0.1, 'eta': 0.1, 'colsample_bytree': 1.0}
RMSE: 13.40118747257878
RMSLE: 0.4652853643060961
R-squared: 0.039336202277775234

Evaluation Metrics

```
[86]: # Evaluate the models
```

```
# Output metrics for comparison
print(f"Linear Regression - RMSLE: {test_rmsle}")
print(f"Decision Tree - RMSLE: {rmsle_test_dt}")
print(f"Random Forest - RMSLE: {rmsle_test_rf}")
print(f"XGBoost - RMSLE: {rmsle_xgb_test}")
for name, metrics in results_rsc.items():
    print(f"RandomizedSearchCV - {name} - RMSLE: {metrics['RMSLE']}")
```

Linear Regression - RMSLE: 0.47250227456368893

Decision Tree - RMSLE: 0.6602510840912299

Random Forest - RMSLE: 0.4658704646094586

XGBoost - RMSLE: 0.4654407429199889

RandomizedSearchCV - Decision Tree - RMSLE: 0.4675707029107024

RandomizedSearchCV - Random Forest - RMSLE: 0.4659722004214911

RandomizedSearchCV - XGBoost - RMSLE: 0.4652853643060961

Save the Best Model

```
[87]: import joblib
```

```
# Save the best model
joblib.dump(best_model, 'kaggle/working/health-insurance-premium-best-model.
    .pkl')
print("Model saved successfully!")
```

Model saved successfully!